

ECE421 - Winter 2022

Assignment 1: Logistic Regression

Faranak Dayyani, Student Number: 1002373674

Part 1: Logistic Regression with Numpy

1. Loss Function and Gradient

The total loss function is as follows:

$$\mathcal{L} = \mathcal{L}_{\text{CE}} + \mathcal{L}_{\text{w}}$$
$$= \frac{1}{N} \sum_{n=1}^N \left[-y^{(n)} \log \hat{y}(\mathbf{x}^{(n)}) - (1 - y^{(n)}) \log (1 - \hat{y}(\mathbf{x}^{(n)})) \right] + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

The Gradient of the loss function with respect to weight (w):

$$\frac{\partial L}{\partial \mathbf{w}} = \frac{1}{N} [(\hat{y} - y) \mathbf{x}^T] + \lambda \mathbf{w}$$

The Gradient of the loss function with respect to bias (b):

$$\frac{\partial L}{\partial b} = \hat{y} - y$$

Code snippet for Loss function:

```
32 # sigmoid function
33 def sigmoid(x,w):
34     z = np.dot(x, w)
35     s = 1/(1 + np.exp(-z))
36
37     return s
38
39
40 def loss(w, b, x, y, reg):
41
42     global LW, L, LCE, LCE_mean, y_hat
43
44     y_hat = sigmoid(x, w)
45
46     w = np.delete(w, -1)
47     w = w.reshape(1, 784)
48     LCE = ((-1*y)*(np.log(y_hat))) - ((1 - y)*np.log(1 - y_hat))
49     LCE_mean = np.mean(LCE) #cross entropy loss value
50
51     w_norm = np.linalg.norm(w) #takes the norm L2
52     LW = (reg/2)*((w_norm)**2) #regularization loss
53     L = LCE_mean + LW #total loss
54
55     return L
```

Figure 1: Loss Function code snippet

Code snippet for Gradient of the Loss function:

```
58 def grad_loss(w, b, x, y, reg):
59
60     global Lgrad_wrt_w, Lgrad_wrt_b
61
62     y_hat = sigmoid(x, w)
63
64     m = len(x)
65     y_result = y_hat - y
66     x_tran = x.T
67     Lgrad_wrt_w = ((1/m)*(np.dot(x_tran, y_result))) + (reg*w)
68     Lgrad_wrt_b = Lgrad_wrt_w[784,0]
69     Lgrad_wrt_w = np.delete(Lgrad_wrt_w,-1)
70     Lgrad_wrt_w = Lgrad_wrt_w.reshape(784,1)
71     return Lgrad_wrt_w, Lgrad_wrt_b
```

Figure 2: Gradient of the Loss Function code snippet

2. Gradient Descent Implementation

```
75 def grad_descent(w, b, x, y, xvalid, yvalid, xtest, ytest, alpha, epochs, reg, error_tol):
76
77     global w_updated, b_updated, ltrain, lvalid, larr_train, xaxis, larr_train_a, larr_valid, larr_valid_a, y_hat_train, y_hat_valid
78     larr_train = 0
79     larr_valid = 0
80     Aarr_train = 0
81     Aarr_valid = 0
82
83     for i in range(epochs):
84
85         # loss data
86         ltrain = loss(w, b, x, y, reg)
87         lvalid = loss(w, b, xvalid, yvalid, reg)
88         ltest = loss(w, b, xtest, ytest, reg)
89
90         # accuracy data
91         y_hat_train = sigmoid(x, w)
92         y_hat_train[y_hat_train >= 0.5] = 1
93         y_hat_train[y_hat_train < 0.5] = 0
94         Atrain = np.mean(y_hat_train == y)
95         Atrain = Atrain*100
96
97         y_hat_valid = sigmoid(xvalid, w)
98         y_hat_valid[y_hat_valid >= 0.5] = 1
99         y_hat_valid[y_hat_valid < 0.5] = 0
100        Avalid = np.mean(y_hat_valid == yvalid)
101        Avalid = Avalid*100
102
103        y_hat_test = sigmoid(xtest, w)
104        y_hat_test[y_hat_test >= 0.5] = 1
105        y_hat_test[y_hat_test < 0.5] = 0
106        Atest = np.mean(y_hat_test == ytest)
107        Atest = Atest*100
108
109        Lgrad_wrt_w, Lgrad_wrt_b = grad_loss(w, b, x, y, reg)
110
111        b = w[784,0]
112        w = np.delete(w,-1)
113        w = w.reshape(784,1)
```

Figure 3: Gradient Descent function code snippet – 1

```

116     w_updated = w - (alpha*Lgrad_wrt_w)
117     b_updated = b - (alpha*Lgrad_wrt_b)
118
119     w_difference = w_updated - w
120     error = np.linalg.norm(w_difference)
121
122     if error < error_tol:
123
124         print("error is less than error tolerance")
125
126         # ---- plotting Loss ----
127         # array for plotting train data
128         Larr_train = np.append(Larr_train,Ltrain)
129         Larr_train_a = Larr_train[1:epochs+1]
130
131         #array for plotting validation data
132         Larr_valid = np.append(Larr_valid,Lvalid)
133         Larr_valid_a = Larr_valid[1:epochs+1]
134
135         xaxis = np.linspace(0, epochs, num=epochs)
136         plt.plot(xaxis,Larr_train_a)
137         plt.plot(xaxis,Larr_valid_a)
138         plt.xlabel('Epochs')
139         plt.ylabel('Loss')
140         plt.title('Loss vs. Number of Epochs - regularization = 0.5')
141         plt.legend(['trainData', 'validData'])
142         plt.grid()
143         plt.show()
144         plt.figure()
145
146         # ---- plotting Accuracy ----
147         #array for plotting train data
148         Aarr_train = np.append(Aarr_train, Atrain)
149         Aarr_train_a = Aarr_train[1:epochs+1]
150
151         #array for plotting validation data
152         Aarr_valid = np.append(Aarr_valid,Avalid)
153         Aarr_valid_a = Aarr_valid[1:epochs+1]

```

Figure 4: Gradient Descent function code snippet – 2

```

155     xaxis = np.linspace(0, epochs, num=epochs)
156     plt.plot(xaxis,Aarr_train_a)
157     plt.plot(xaxis,Aarr_valid_a)
158     plt.xlabel('Epochs')
159     plt.ylabel('Accuracy (%)')
160     plt.title('Accuracy vs. Number of Epochs - regularization = 0.5')
161     plt.legend(['trainData','validData'])
162     plt.grid()
163     plt.show()
164
165     print("error is less than error_tol")
166     break;
167
168     w = w_updated
169     b = b_updated
170     w = np.append(w,b)
171     w = w.reshape(785,1)
172
173     # ---- Loss arrays ----
174     # array for plotting train data
175     Larr_train = np.append(Larr_train,Ltrain)
176     Larr_train_a = Larr_train[1:epochs+1]
177
178     #array for plotting validation data
179     Larr_valid = np.append(Larr_valid,Lvalid)
180     Larr_valid_a = Larr_valid[1:epochs+1]
181
182     # ---- Accuracy arrays ----
183     #array for plotting train data
184     Aarr_train = np.append(Aarr_train, Atrain)
185     Aarr_train_a = Aarr_train[1:epochs+1]
186
187     #array for plotting validation data
188     Aarr_valid = np.append(Aarr_valid,Avalid)
189     Aarr_valid_a = Aarr_valid[1:epochs+1]
190

```

Figure 5: Gradient Descent function code snippet – 3

```

192 # ---- plotting Loss ----
193 xaxis = np.linspace(0, epochs, num=epochs)
194 plt.plot(xaxis,Larr_train_a)
195 plt.plot(xaxis,Larr_valid_a)
196 plt.xlabel('Epochs')
197 plt.ylabel('Loss')
198 plt.title('Loss vs. Number of Epochs - regularization = 0.5')
199 plt.legend(['trainData','validData'])
200 plt.grid()
201 plt.show()
202
203 plt.figure()
204
205 # ---- plotting Accuracy ----
206 xaxis = np.linspace(0, epochs, num=epochs)
207 plt.plot(xaxis,Aarr_train_a)
208 plt.plot(xaxis,Aarr_valid_a)
209 plt.xlabel('Epochs')
210 plt.ylabel('Accuracy (%)')
211 plt.title('Accuracy vs. Number of Epochs - regularization = 0.5')
212 plt.legend(['trainData','validData'])
213 plt.grid()
214 plt.show()
215
216 print("Training Loss:",Ltrain)
217 print("Validation Loss:",Lvalid)
218 print("Testing Loss:",Ltest)
219 print("Training Accuracy:",Atrain)
220 print("Validation Accuracy:",Avalid)
221 print("Testing Accuracy:",Atest)
222
223
224 return w_updated, b_updated

```

Figure 6: Gradient Descent function code snippet – 4

3. Tuning the Learning Rate

Table 1: The table below shows the loss and accuracy values for different learning rate (α) values:

Data Type	$\alpha = 0.005$	$\alpha = 0.001$	$\alpha = 0.0001$
Training Loss	0.054409	0.073337	0.158731
Validation Loss	0.060161	0.083721	0.177008
Testing Loss	0.084959	0.086521	0.158022
Training Accuracy	98.31428	97.8	97.39999
Validation Accuracy	98.0	98.0	97.0
Testing Accuracy	97.93103	97.24137	97.24137




Figure 7: Loss vs. Number of Epochs for $\alpha = 0.005$

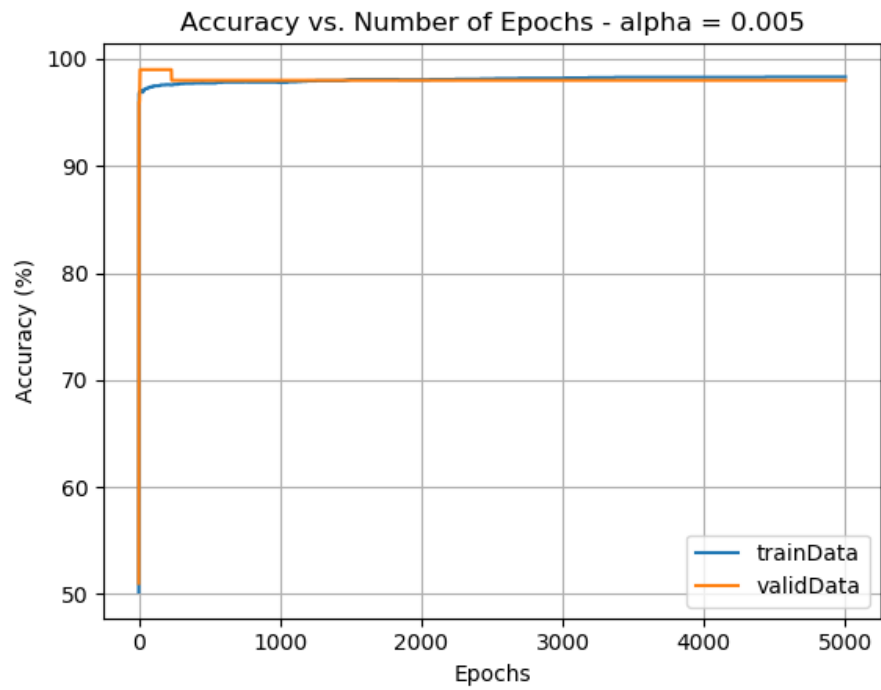


Figure 8: Accuracy vs. Number of Epochs for $\alpha = 0.005$




Figure 9: Loss vs. Number of Epochs for $\alpha = 0.001$




Figure 10: Accuracy vs. Number of Epochs for $\alpha = 0.001$

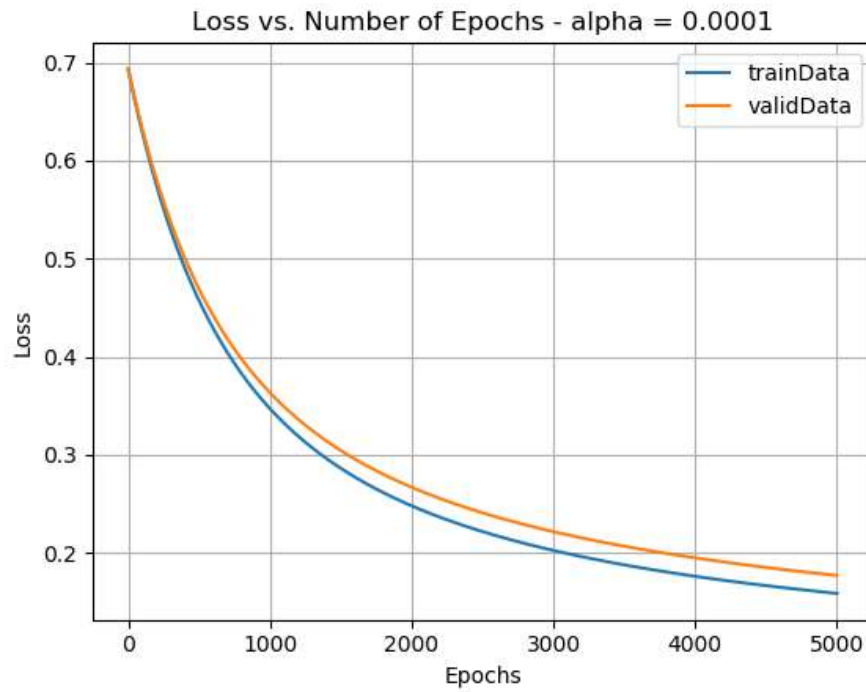


Figure 11: Loss vs. Number of Epochs for $\alpha = 0.0001$




Figure 12: Accuracy vs. Number of Epochs for $\alpha = 0.0001$

Figure 7-12 above show the loss and accuracy curves for training and validation sets for different learning rate (α) values. Table 1 above displays the loss and accuracy values for training, validation and testing sets for different learning rate (α) values. According to the figures 7-12 and

table 1, it can be concluded that the best learning rate is 0.005. This can be supported by the lowest testing loss (0.084959) and highest testing accuracy (97.93103%) observed in table 1. This learning rate also provides the highest accuracy and lowest loss on both of the training and validation sets. The graphs illustrate that the learning rate of 0.005 converges faster than the other two.

In addition, when decreasing the learning rate value to 0.0001, the model has a poor loss compared to the other two values. This is due to the smaller step size which causes it to take more time to reach the global minimum. The smaller the learning rate, the more epochs required to reach the global minimum.

4. Generalization

Table 2: The table below shows the loss and accuracy values for different values of regularization parameter (λ):

Data Type	$\lambda = 0.001$	$\lambda = 0.1$	$\lambda = 0.5$
Training Loss	0.055624	0.114683	0.198496
Validation Loss	0.061497	0.126721	0.216019
Testing Loss	0.085774	0.127311	0.201335
Training Accuracy	98.31429	98.08571	97.74286
Validation Accuracy	98.0	98.0	98.0
Testing Accuracy	97.93103	97.93103	97.93103

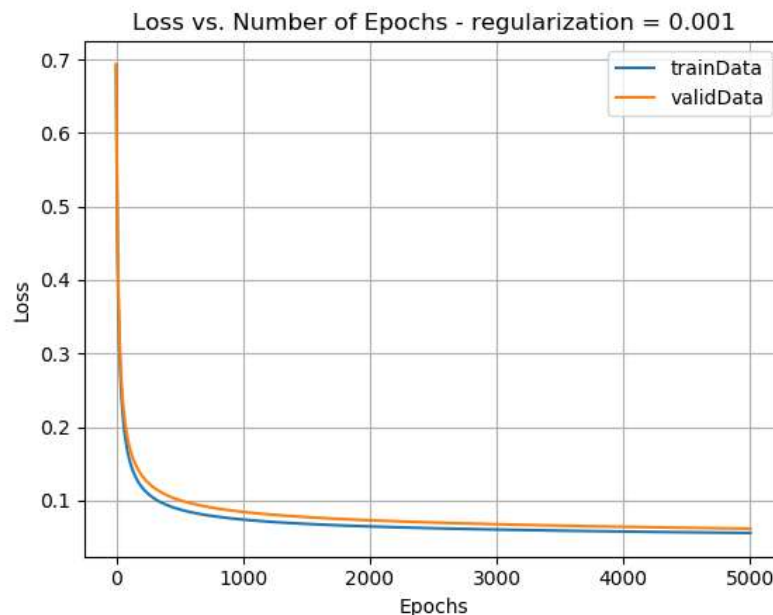


Figure 13: Loss vs. Number of Epochs for $\lambda = 0.001$




Figure 14: Accuracy vs. Number of Epochs for $\lambda = 0.001$




Figure 15: Loss vs. Number of Epochs for $\lambda = 0.1$




Figure 16: Accuracy vs. Number of Epochs for $\lambda = 0.1$

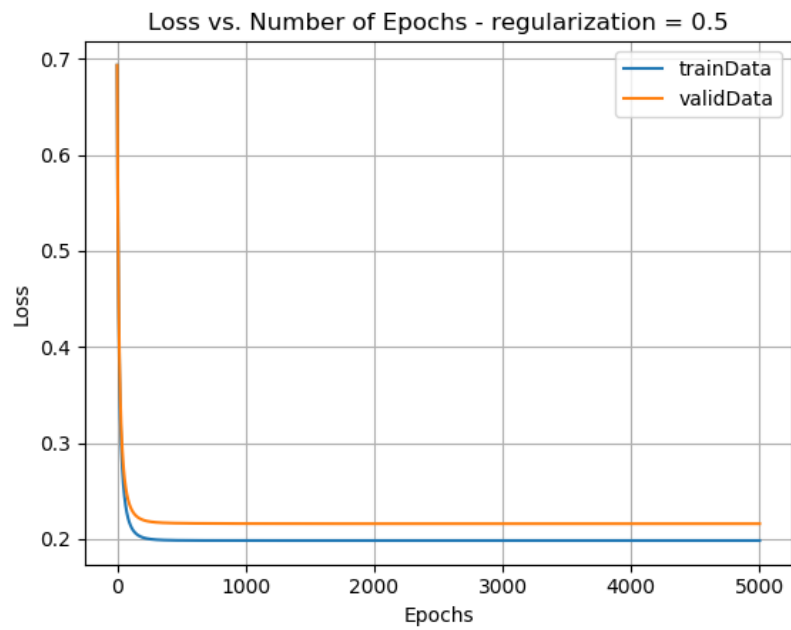


Figure 17: Loss vs. Number of Epochs for $\lambda = 0.5$

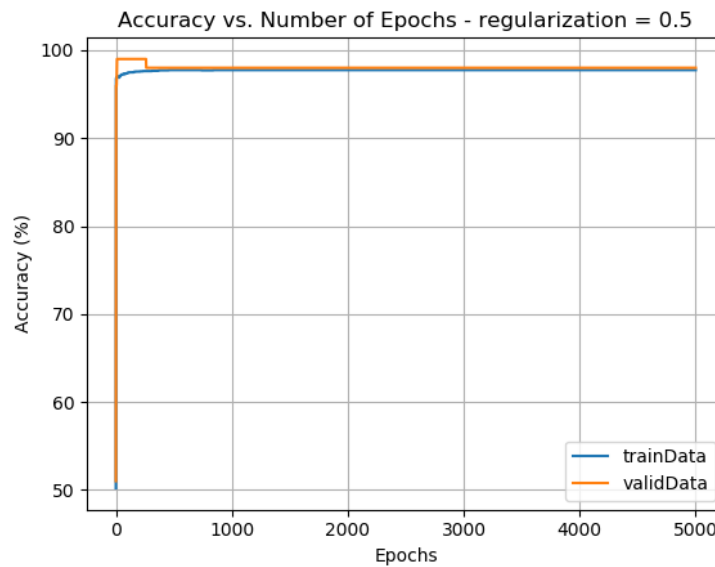


Figure 18: Accuracy vs. Number of Epochs for $\lambda = 0.5$

Figure 13-18 above show the loss and accuracy curves for training and validation sets for different values of the regularization (λ) parameter. Table 2 above displays the loss and accuracy values for training, validation and testing sets for different regularization (λ) parameter values. Regularization is used to limit the growth of weight which ensures the curve doesn't fit too well on the training data. According to table 2, it can be seen that as we increase the value of λ , the training loss value also increases. However, lower values of λ contribute less to the loss value in order to prevent overfitting.

According to figures 13-18 and table 2, it can be concluded that $\lambda = 0.5$ is the best regularization parameter value as the model avoids overfitting on the training data. The convergence rate is higher for $\lambda = 0.5$ compared to the other two values. In addition, the test accuracy (97.93103%) is greater than the training accuracy (97.74286%) for $\lambda = 0.5$ which supports the above statement.

Part 2: Logistic Regression in TensorFlow

1. Building the Computational Graph

```
[ ] def buildGraph(xdata, ydata, alpha, ADAM, beta_1=0.9, beta_2=0.999, epsilon=1e-08):
    global loss
    tf.set_random_seed(20)
    xdata_d = xdata.shape[1]

    # ----- Question 1 - Part a -----
    #initializing weight and bias tensors
    w = tf.truncated_normal_initializer(mean=0, stddev=0.5)
    w = tf.get_variable("w", (xdata_d, 1), initializer=w)
    b = tf.constant_initializer(0.0)
    b = tf.get_variable("b", (1, ), initializer=b)

    # ----- Question 1 - Part b -----
    # placeholders for data, label and reg
    x = tf.placeholder(tf.float32, shape=(None, xdata_d))
    y = tf.placeholder(tf.float32, shape=(None, 1))
    reg = tf.placeholder(tf.float32)

    # ----- Question 1 - Part c -----
    # Loss function
    z = tf.add(tf.matmul(x, w), b)
    y_hat = tf.nn.sigmoid(z, name="y_hat")
    LCE = tf.nn.sigmoid_cross_entropy_with_logits(labels=y, logits=z)
    LCE = tf.reduce_mean(LCE) #cross entropy loss
    LW = (reg/2) * tf.reduce_sum(tf.square(w)) #regularization term loss
    loss = tf.add(LCE, LW) #total loss

    # ----- Question 1 - Part d -----
    # optimizer
    if ADAM==True:
        optimizer = tf.train.AdamOptimizer(alpha, beta_1, beta_2, epsilon).minimize(loss)
    else:
        optimizer = tf.train.GradientDescentOptimizer(alpha).minimize(loss)

    return x, w, b, y_hat, y, loss, optimizer, reg
```

Figure 19: Code snippet of buildGraph function

2. Implementing Stochastic Gradient Descent

```
def SGD(batch_size, epochs, alpha, reg_2, ADAM, beta_1=0.9, beta_2=0.999, epsilon=1e-08):
    global trainData, validData, trainTarget, validTarget, testData, testTarget, loss_2, loss

    N = trainData.shape[0]
    total_batches = int(N/batch_size)
    iterations = total_batches * epochs

    x, w, b, y_hat, y, loss, optimizer, reg = buildGraph(trainData, trainTarget, alpha, ADAM, beta_1, beta_2, epsilon)

    global_init = tf.global_variables_initializer()
    L_train = []
    L_valid = []
    A_train = []
    A_valid = []
    L_test = []
    A_test = []

    with tf.Session() as sess:
        sess.run(global_init)
        for i in range(iterations):
            if (i+1)%total_batches == 0:
                shuffle_index = np.random.choice(N, N, replace=False)
                trainData, trainTarget = trainData[shuffle_index], trainTarget[shuffle_index]
                trainBatch, ybatch = trainData[:batch_size], trainTarget[:batch_size]

                _, w_2, b_2, loss_2 = sess.run([optimizer, w, b, loss], feed_dict={x:trainBatch, y:ybatch, reg:reg_2})
                trainData = np.roll(trainData, batch_size, axis=0)
                trainTarget = np.roll(trainTarget, batch_size, axis=0)
```

Figure 20: Code snippet of SGD function – 1

```
            if (i+1)%total_batches == 0:
                L_train_2, L_valid_2, A_train_2, A_valid_2, L_test_2, A_test_2 = AccValue(w_2, b_2, trainData, trainTarget, reg_2)
                L_train.append(L_train_2)
                L_valid.append(L_valid_2)
                A_train.append(A_train_2)
                A_valid.append(A_valid_2)
                L_test.append(L_test_2)
                A_test.append(A_test_2)

            w_updated, b_updated = sess.run([w, b])

    return w_updated, b_updated, L_train, L_valid, A_train, A_valid, L_test, A_test
```

Figure 21: Code snippet of SGD function – 2

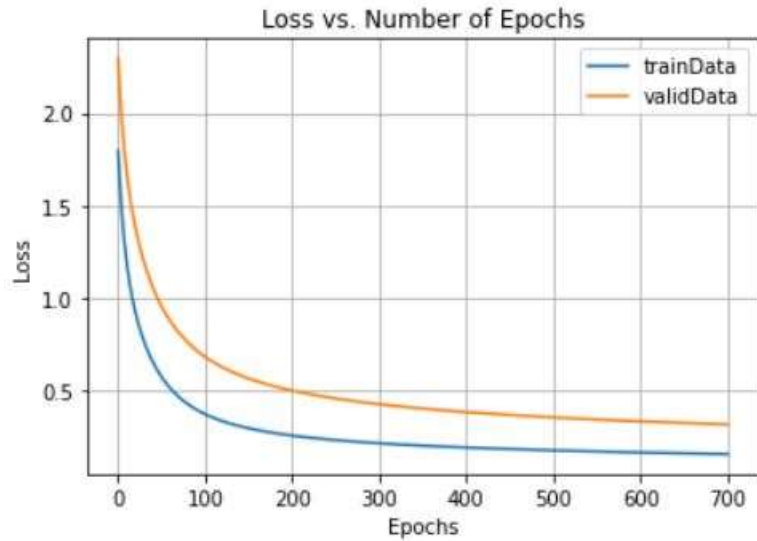


Figure 22: Loss vs. Number of Epochs for SGD

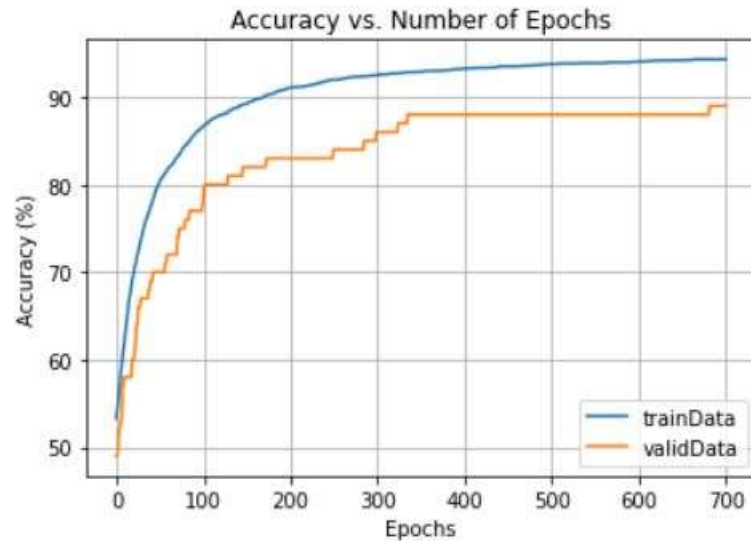


Figure 23: Accuracy vs. Number of Epochs for SGD

3. Batch Size Investigation

Table 3: The table below shows the loss and accuracy values for different batch sizes:

Data Type	Batch size = 100	Batch size = 700	Batch size = 1750
Training Loss	0.005035	0.026014	0.051225
Validation Loss	0.136783	0.110864	0.144830
Testing Loss	0.202145	0.123003	0.138895
Training Accuracy	99.97143	99.02857	97.94286
Validation Accuracy	97.0	96.0	95.0
Testing Accuracy	97.24138	98.62069	97.24138

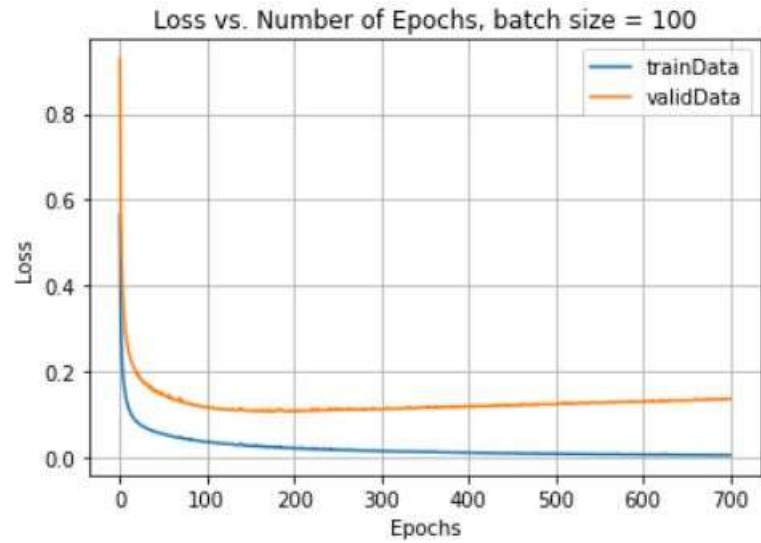


Figure 24: Loss vs. Number of Epochs for SGD – batch size = 100

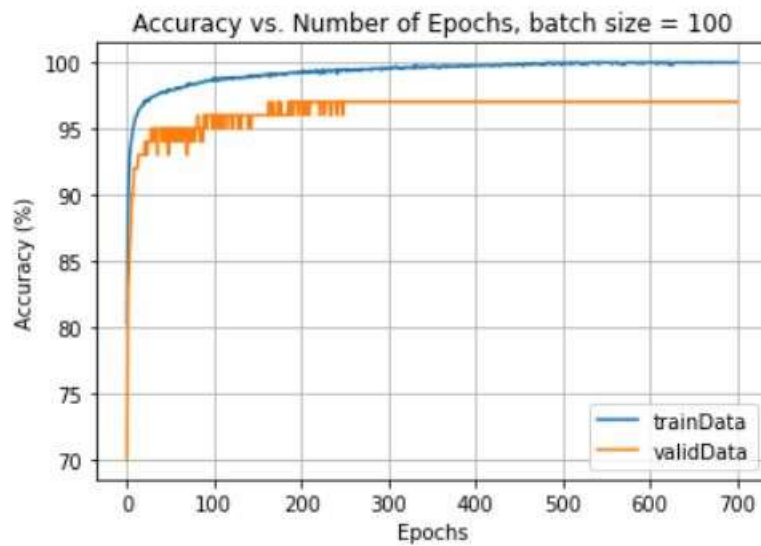


Figure 25: Accuracy vs. Number of Epochs for SGD – batch size = 100




Figure 26: Loss vs. Number of Epochs for SGD – batch size = 700




Figure 27: Accuracy vs. Number of Epochs for SGD – batch size = 700

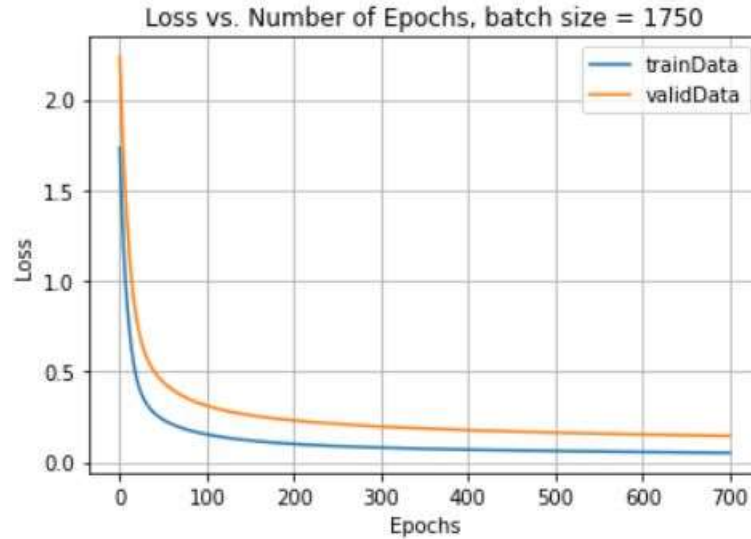


Figure 28: Loss vs. Number of Epochs for SGD – batch size = 1750




Figure 29: Accuracy vs. Number of Epochs for SGD – batch size = 1750

Figure 24-29 above show the loss and accuracy curves for training and validation sets for different batch sizes. Table 3 above displays the loss and accuracy values for training, validation and testing sets for different batch size values. According to the figures, it can be stated that the convergence rate of the training data for batch size = 100 is slower in comparison with the higher batch sizes. However, the convergence rate for the validation data of batch size = 100 does not reach the global minimum. As seen in figure 27 and 29, the accuracy curves for higher batch sizes are fairly stable. The best batch size value is 100 since it has the highest validation accuracy (97.0%). The highest testing accuracy is for batch size = 700 but when looking at training and validation accuracy values coupled with the testing accuracy, the model performs best with batch size = 100. The small batches have a regularizing effect by adding noise to gradient estimations. This leads to better generalization by helping the model to come out of the local minima. To

conclude, as the batch size increases, the effect of regularization decreases which leads to lower test accuracy.

4. Hyperparameter Investigation

Table 4: The table below shows the loss and accuracy values for different Adam hyperparameters:

Data Type	$\beta_1 = 0.95$	$\beta_1 = 0.99$	$\beta_2 = 0.99$	$\beta_2 = 0.9999$	$\varepsilon = 1e - 09$	$\varepsilon = 1e - 4$
Training Loss	0.019307	0.019437	0.011871	0.026001	0.019329	0.019729
Validation Loss	0.109335	0.110801	0.116254	0.1121769	0.108994	0.108917
Testing Loss	0.129573	0.129362	0.146202	0.1229372	0.129105	0.128440
Training Accuracy	99.34574	99.34286	99.6	99.05714	99.34286	99.34286
Validation Accuracy	97.0	97.0	97.0	96.0	97.0	97.0
Testing Accuracy	98.62189	98.62069	97.24138	96.62069	98.62179	98.62069

❖ Note: the default hyperparameters that were used for Adam are:

- $\beta_1 = 0.9$
- $\beta_2 = 0.999$
- $\varepsilon = 1e - 08$

Table 4 above displays the loss and accuracy values for training, validation and testing sets for different values of Adam hyperparameters. β_1 and β_2 are used for speeding up the gradient descent in Adam.

- a) β_1 is the exponential decay rate for the first moment estimate. As seen in table 4 above, $\beta_1 = 0.95$ has higher accuracy values compared to $\beta_1 = 0.99$. The lower the value of β_1 , the more importance is given to the historical gradient values for updating the first moment estimate which explains the reason behind $\beta_1 = 0.95$ resulting in higher accuracy values compared to $\beta_1 = 0.99$. To conclude, $\beta_1 = 0.95$ is a better pick due to its higher test accuracy value (98.62189%).
- b) β_2 is the exponential decay rate for the second moment estimate. As seen in table 4 above, $\beta_2 = 0.99$ has higher accuracy values compared to $\beta_2 = 0.9999$. The higher the value of β_2 , the more weight is given to the square of historical gradient for updating the second moment estimate which explain the reason behind $\beta_2 = 0.99$ resulting in higher accuracy values compared to $\beta_2 = 0.9999$. To conclude, $\beta_2 = 0.99$ is a more ideal pick due to its higher test accuracy value (97.24138%).
- c) ε value is used to prevent a division by zero in the implementation of the model. As seen in table 4 above, $\varepsilon = 1e - 09$ has slightly higher accuracy values compared to $\varepsilon = 1e - 4$ while being very close. To conclude, $\varepsilon = 1e - 09$ is a more ideal pick due to its slightly higher test accuracy value (98.62179%).

5. Comparison against Batch GD

Table 5: The table below shows the loss and accuracy values for Stochastic Gradient Descent with Adam and Batch Gradient Descent:

Data Type	ADAM	Batch Gradient Descent
Training Loss	0.012142	0.073337
Validation Loss	0.102934	0.083722
Testing Loss	0.157186	0.0865216
Training Accuracy	99.62857	97.8
Validation Accuracy	97.0	98.0
Testing Accuracy	97.93103	97.24138

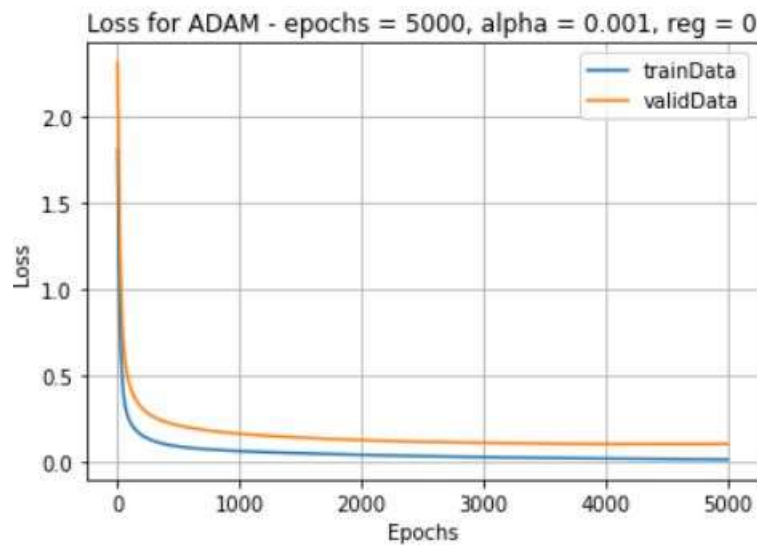


Figure 30: Loss vs. Number of Epochs for SGD with Adam

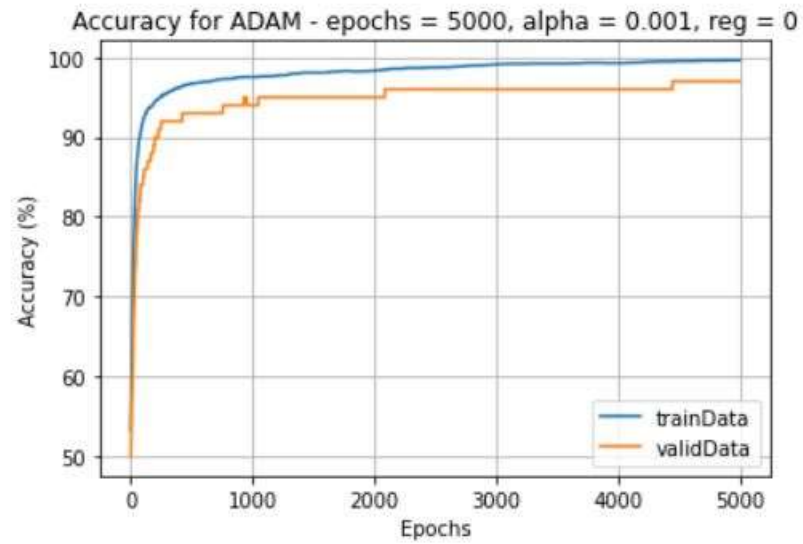


Figure 31: Accuracy vs. Number of Epochs for SGD with Adam

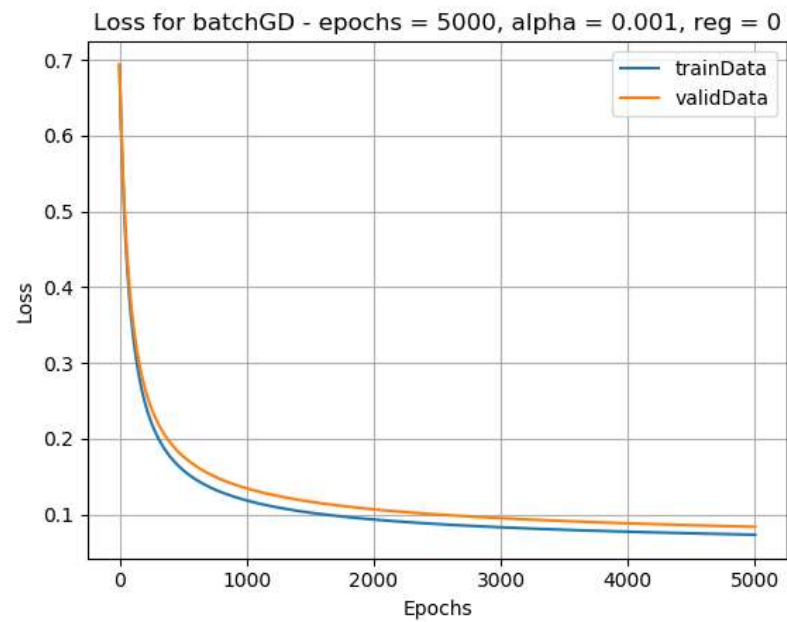


Figure 32: Loss vs. Number of Epochs for Batch Gradient Descent

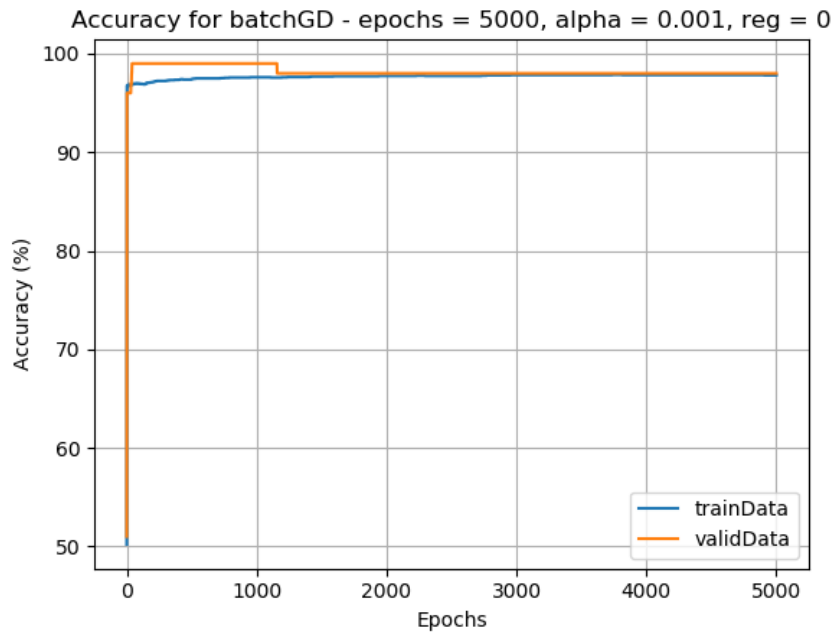


Figure 33: Accuracy vs. Number of Epochs for Batch Gradient Descent

Figures 30-33 above shows the loss and accuracy curves for training and validation sets for Stochastic Gradient Descent (SGD) algorithm with Adam and Batch Gradient Descent (GD) algorithm. Table 5 above illustrates the loss and accuracy values for SGD with Adam and Batch GD with the following hyperparameters:

- Epochs = 5000
- $\lambda = 0$
- $\alpha = 0.001$

The hyperparameters were kept the same for both in order to have a fair assessment of the overall performance of both algorithms. The overall performance of SGD algorithm is better than the Batch GD algorithm due to its use of mini-batches.

According to figures 30 and 32 above, it can be seen that the final loss for SGD is lower than that of Batch GD. Furthermore, the rate of convergence is higher in SGD.

According to figures 31 and 33 above, it can be seen that the final accuracy of SGD is higher than the Batch GD.

To conclude, SGD algorithm with Adam has a better performance in comparison with Batch GD. One of the reasons for this can include that the Adam optimizer aids in a faster convergence rate.